

KUMAR GYANAM

Chief AI Architect | AI Strategy Leader

📍 Gurgaon | 📞 9953682525 | 📩 gyanamc@gmail.com | 🔗 [LinkedIn Profile](#) | 🌐 [Live Portfolio](#): gyanam.store

STOP READING. START INTERVIEWING. “I built an Agentic AI bot to deprecate my own traditional resume. Instead of reading, if you would like to talk to my resume, then please click”  [CLICK HERE](#)

EXECUTIVE SUMMARY

Transformational AI Leader with **~20 years of experience** architecting enterprise-scale AI systems across fintech, healthcare, and media. Proven expert in building and deploying **sovereign, on-premise GenAI infrastructure** within highly regulated environments. Trusted advisor to C-level stakeholders with board-approved AI roadmaps, managing full-cycle AI transformation, multi-million dollar budgets, and vendor ecosystems. **Guest Faculty at IIM and ISB.**

TECHNICAL ARCHITECTURE EXPERTISE

- **Agentic AI & Orchestration:** LangGraph (Stateful Agents), LangChain, n8n, Agentic RAG Workflows.
 - **On-Premise Generative AI Stack:** Ray Serve (Distributed Serving), vLLM (Inference Optimization), ELK Stack (Vector Search), Open Source LLMs (Llama-3, Mistral, OSS).
 - **MLOps & Observability:** Ragas (RAG Evaluation), MLflow (Experiment Tracking & Registry), Airflow (Pipeline Orchestration), Prometheus/Grafana.
 - **Cloud & Platforms:** Railway.app (Expert—30+ deployments), AWS, Firebase, Docker, Kubernetes (K8s).
 - **Data & Automation:** n8n.io (Expert—30+ workflows), Python (scikit-learn, TensorFlow), FAST APIs, SQL/NoSQL.
-

PROFESSIONAL EXPERIENCE

SBI Card | Deputy Vice President - Conversational AI

April 2023 – Present

Strategic AI Infrastructure & Platform Leadership

- **Sovereign On-Premise Agentic RAG Platform:** Architected and operationalised a secure, air-gapped Retrieval-Augmented Generation (RAG) ecosystem, ensuring 100% data sovereignty and financial-grade security.

- **Orchestration & Logic:** Leveraged **LangChain** and **LangGraph** to build complex, stateful agentic workflows, enabling multi-turn reasoning and tool-integrated AI capabilities.
 - **High-Performance Stack:** Integrated **45+ enterprise components** including **Ray Serve** for distributed serving and **vLLM** for high-throughput inference, managed via **Airflow** for robust data pipeline orchestration.
 - **Infrastructure Depth:** Deployed open-source LLMs (**Mistral**, **Llama-3**) using **ELK (Elasticsearch)** as a **high-scale Vector Database** for efficient semantic retrieval.
 - **Continuous Evaluation:** Established a rigorous evaluation loop using **Ragas** (for metric-based assessment of faithfulness/relevancy) and **MLflow** for end-to-end experiment tracking and versioning.
- **Multi-Modal Strategic Implementations:** Utilized the on-premise RAG engine to deliver high-impact business applications:
 1. **Credit Card Recommendation Engine:** Real-time discovery engine tailored to customer financial profiles and spending behavior.
 2. **Generative Agent-Assist:** Real-time reply generation for live chat agents, improving response accuracy and reducing average handle time (AHT).
 3. **SBI Card Internal GPT:** A secure, internal-only portal for employees to interact with massive knowledge bases and internal documentation.
 4. **Self-Healing & HITL Architecture:** Engineered a **Human-in-the-Loop (HITL)** framework to allow real-time correction of model misses or biases, ensuring reliability.
- **Platform Scale & Governance:** Scaled the conversational platform to support **22M+** users and **21M+ monthly transactions** with a **99.9% uptime SLA**.
 - Authored the board-approved AI roadmap, driving a **400% growth in bill payments** via AI channels and leading enterprise-wide AI procurement and vendor evaluation.

Brane Group | Analytics Delivery Head

January 2022 – April 2023

- Architected ML-powered solutions for Healthcare (ambulance dispatch optimisation), Financial Risk (fraud detection), and Ad-tech (campaign ROI optimisation).
- Led a cross-functional team of data scientists and engineers to deliver production-ready AI models for global clients.

Pixstory | Head of Technology Architecture*January 2019 – January 2022*

- Built a custom ML-based recommendation engine from scratch, resulting in a **180% improvement in user engagement.**
- Led technical product strategy and executed go-to-market strategies supporting 100K+ user acquisition.
- Managed the full-stack development of semantic analysis and automated tagging workflows for 1M+ posts.

The Indian Express | Chief Manager - Business Intelligence*July 2009 – January 2019*

- Transformed subscription architecture and built predictive analytics for collection management.
- Reduced user detractors by **50%** through text mining and continuous consumer research.

Axis Bank | Sales Manager*August 2008 - July 2009*

Managed distributed sales operation achieving Rs 200 Cr+ disbursement targets through process optimisation and team leadership.

ICICI Bank (Sales Manager) |*June 2006 - August 2008:*

Built franchise network across multiple districts managing Rs 60 Cr+ lending portfolio and regional expansion strategy.

EDUCATION & ACADEMIC LEADERSHIP

- **Academic Leadership:** Faculty at **IIM-Indore** and **ISB** for AI Strategy, Technology Leadership, and AI Governance.
 - **Indian Institute of Technology (IIT), Madras:** Advanced Diploma in Programming (2022–2025, Dropped Out).
 - **Indian Institute of Management (IIM), Calcutta:** AI-Powered Marketing (2021–2022).
 - **Institute of Rural Management (IRM):** MBA in Marketing & Strategic Management (2004–2006).
 - **Ranchi University:** B.Sc. in Mathematics & Statistics (2001–2004).
-

KEY ACHIEVEMENTS & STRATEGIC IMPACT

- **Production Excellence:** Maintains a live portfolio at **gyanam.store** featuring 30+ production AI systems.

- **Operational Transformation:** Achieved **75+ FTE cost savings** through live chat integration and AI-driven automation.
- **Thought Leadership:** Recognised expert in **Railway.app** and **n8n.io** for enterprise-scale AI automation.
- **Governance:** Expert in fintech AI governance, regulatory compliance, and large-scale system security audits.